AgentShield: FF 180 Major Project Synopsis

1. Environment

AgentShield considers environmental sustainability by directly reducing wasteful GPU and cloud computing overhead caused by autonomous AI agents. When AI agents fail a task, they frequently enter 'hallucination loops,' continuously firing redundant API calls to external tools.

AgentShield implements a sliding-window rate limiter and quarantine protocol that detects and halts these runaway loops, significantly reducing API traffic, GPU token generation, and cloud energy consumption. Furthermore, AgentShield uses a deterministic, rule-based security engine rather than a heavy Machine Learning model, resulting in an ultra-low carbon footprint.

Environmental Impact & Compute Optimization Example:

- API Calls: 100+ calls/min (Without Proxy) vs Capped at 10 calls/5sec (With AgentShield)
- Token Waste: ~50,000+ wasted tokens/min vs Cut off instantly (Zero waste)
- Security Compute: Heavy LLM API calls vs Local Regex/Heuristics (Sub-2ms)
- Memory Usage: Heavy vs < 50 MB RAM

2. Sustainability

Sustainability addresses both technical and operational longevity. AgentShield is designed as a modular, protocol-agnostic system that supports long-term technical sustainability. Because it strictly adheres to the official JSON-RPC 2.0 standard for the Model Context Protocol (MCP), it guarantees compatibility with future AI models and enterprise tools without requiring codebase rewrites.

AgentShield Project-Specific Considerations:

- Modular: Separated logically into Gateway, Risk Engine, Sanitizer, and HITL Manager.
- Extensible: New attack vectors and PII patterns can be added by simply appending new Regex rules without touching core routing logic.
- Vendor-Agnostic: Compatible with any future AI client (e.g., Claude, GPT) or MCP server.

Example of Extensibility:

The PII Sanitizer currently redacts: Aadhaar -> PAN -> AWS API Keys.
Later, new compliance requirements can be added: Credit Cards -> Health Records -> Internal IP Addresses simply by updating src/sanitizer/patterns.py.

3. Safety and Security

AgentShield is fundamentally a cybersecurity project. It ensures the safety of enterprise infrastructure by providing multi-tiered runtime guardrails between unpredictable AI agents and sensitive corporate tools.

AgentShield Project-Specific Safety Considerations:

- Prompt Injection Defense: Blocks malicious jailbreaks (e.g., <system_override>, DAN modes) before they reach the execution environment.
- Destructive Action Prevention: Uses dynamic wildcards to flag dangerous tools (e.g., containing 'delete', 'drop', 'execute').
- Human-in-the-Loop (HITL): Pauses high-risk commands in server memory and waits for explicit human administrator approval via the dashboard.
- Fail-Closed Architecture: If an administrator does not approve a HITL request within 60 seconds, the request is

automatically destroyed.

Example:

If an AI agent attempts to call the 'github_delete_repo' tool, AgentShield does not automatically reject or accept it. Instead, it assigns a HIGH RISK score, holds the JSON payload asynchronously, and pushes an alert to the Admin Dashboard for manual verification.

AgentShield: FF 180 Major Project Synopsis

4. Ethical Considerations

AgentShield inherently enforces Ethical AI practices by guaranteeing data privacy, system transparency, and human oversight. Because MCP allows AI to read local enterprise files and databases, there is a massive risk of sensitive corporate data being sent to external LLM providers.

Important Ethical Considerations Addressed:

- Data Privacy (DLP): The Outbound PII Sanitizer recursively scrubs returning payloads, replacing sensitive personal information with '[REDACTED]' before it is transmitted back to the AI.

- Accountability & Transparency: Every single intercepted request, risk score, and action is logged asynchronously to agentshield_audit.db. If an AI makes a mistake, there is a perfect, tamper-proof audit trail.

- Human Agency: By implementing HITL, AgentShield ensures that AI cannot take autonomous, destructive actions without human consent, aligning with global Responsible AI guidelines.

5. Cost and Economic Feasibility

The economic feasibility of AgentShield is exceptionally high. It is built entirely on open-source technologies, incurring zero commercial licensing fees, and its deterministic risk engine runs locally, eliminating the need to pay for external LLM API calls for security checks.

Potential Economic Benefits: The system prevents catastrophic financial losses associated with AI-driven data breaches, accidental database deletions, and runaway API token costs.

Cost Assessment (Without AgentShield vs. With AgentShield):

- Security Licensing: Expensive API gateways vs Rs. 0 (100% Open Source)

- AI Token Waste: High (Prone to infinite loops) vs Negligible (Quarantines rogue agents)

- Data Breach Risk: High (Raw data sent to LLM providers) vs Low (Automated PII Redaction)

- Incident Auditing: Manual log parsing vs Automated SQLite JSONB tracking

- Infrastructure Req.: High (if using ML) vs Minimal (<50MB RAM, sub-2ms latency)

6. SDG Project Mapping

Goal 9: Industry, Innovation, and Infrastructure

Explanation: By building a resilient, low-latency security proxy, AgentShield allows enterprises to safely adopt cutting-edge AI technologies without risking their existing infrastructure. It fosters innovation by removing the security roadblocks associated with autonomous AI deployment.

Goal 16: Peace, Justice, and Strong Institutions

Explanation: AgentShield enforces accountability and transparency in AI systems. Through its exhaustive audit logging and Human-in-the-Loop approval gates, it ensures that corporate institutions maintain strong, governed oversight over how AI interacts with their private data, promoting digital trust.